

Solving Problems by Uninformed Searching

Tian-Li Yu

Taiwan Evolutionary Intelligence Laboratory (TEIL)
Department of Electrical Engineering
National Taiwan University
tianliyu@ntu.edu.tw

Readings: AIMA Sections 3.1~3.4

Outline

- 1 Problem-Solving Agents
- 2 Problem Formulation
- 3 Search on Trees and Graphs
- 4 Uninformed Search
 - Breadth-First
 - Uniform-Cost
 - Depth-First
 - Depth-Limited
 - Iterative Deepening

Problem-Solving Agents

- A simple problem-solving agent **formulates a goal and a problem**, searches for **a sequence of actions** that solves the problem, and then **execute the actions** one by one.

SIMPLE-PROBLEM-SOLVING-AGENT(*percept*)

```

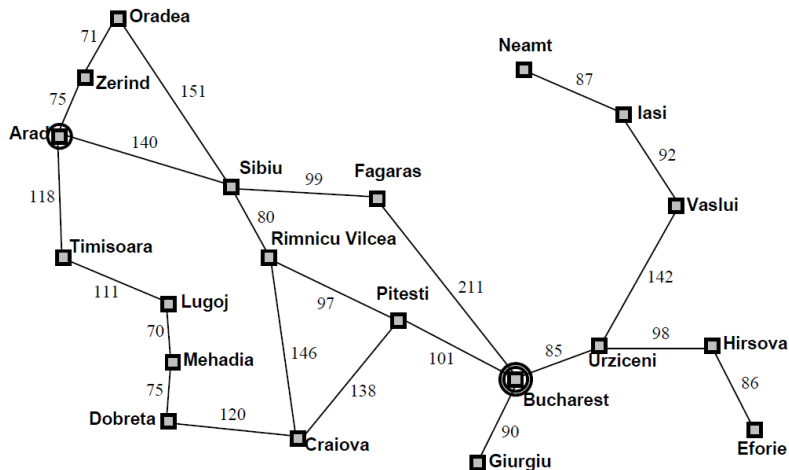
1  state = UPDATE-STATE(state, percept)
2  if seq == empty
3      goal = FORMULATE-GOAL(state)
4      problem = FORMULATE-PROBLEM(state, goal)
5      seq = SEARCH(problem)
6      if seq == failure
7          return NIL
8  action = FIRST(seq)
9  seq = REST(seq)
10 return action

```

Example: Romania

- On holiday in Romania; currently in Arad.
- Flight leaves tomorrow from Bucharest.
- **Formulate goal**
 - Be in Bucharest.
- **Formulate problem**
 - **States**: various cities
 - **Actions**: fly between cities
- **Find solution**
 - Sequence of cities, e.g., Arad, Sibiu, Fagaras, Bucharest

Example: Romania



Problem Formulation

- A **problem** is defined by five components

- 1 Initial state: $ln(Arad)$

- 2 Actions:

$ACTION(ln(Arad)) = \{Go(Sibiu), Go(Timisoara), Go(Zerind)\}$

- 3 Transition model $RESULT(s, a)$:

- $RESULT(ln(Arad), Go(Zerind)) = ln(Zerind)$.

Successor $S(s)$: states reachable by a single action.

- $S(s) = \{s' \mid \forall a \in ACTION(s), s' = RESULT(s, a)\}$

- 4 Goal test: $\{ln(Bucharest)\}$

- 5 Path cost (additive)


- Sum of distances, number of actions executed, etc.
- $c(s, a, s')$ is the **step cost** of taking action a in state s to reach state s' , assumed to be ≥ 0

- A **solution** is a sequence of actions leading from the initial state to a goal state.

Abstraction

- Real world is absurdly complex
 - State space must be **abstracted** for problem solving.
- **(Abstract)** state = subset of real states
- **(Abstract)** action = complex combination of real actions
 - **Go(Zerind)** represents a complex set of possible routes, detours, rest stops, etc.
- For guaranteed realizability, **any** real state “in Arad” must get to **some** real state “in Zerind”
- **(Abstract)** solution = set of real paths that are solutions in the real world
- Each abstract action should be “easier” than the original problem!

Example: Vacuum World State Space Graph



- ① **Initial state:** Any one of the above states. (ignore dirt **amounts** etc.)
- ② **Actions:** Left, Right, Suck, NoOp
- ③ **Transition model:** The above figure.
- ④ **Goal test:** no dirt
- ⑤ **Path cost:** 1 per action (0 for NoOp)

Example: The 8-puzzle

7	2	4
5		6
8	3	1

Start State


1	2	3
4	5	6
7	8	

Goal State

- ① **Initial state:** The left figure.
- ② **Actions:** Move blank left, right, up, down.
- ③ **Transition model:** Common sense.
- ④ **Goal test:** = The right figure.
- ⑤ **Path cost:** One per move.


Note: Sliding-block puzzle is \mathcal{NP} -hard.

Example: 8-Queen Puzzle



- ① **Initial state:** No queen on the board.
- ② **Actions:** Add a queen on the board where the square is empty.
- ③ **Transition model:** Returns the board with a queen added to the specified square.
- ④ **Goal test:** 8 queens are on the board, none attacked.
- ⑤ **Path cost:** Number of trials.

Example: 8-Queen Puzzle



- **States:** Any 0~8 queens on the board.
 - State space: $C_0^{64} + C_1^{64} + C_2^{64} + \dots + C_8^{64} \simeq 5.1 \times 10^9$
 - Solution space: $64 \cdot 63 \dots 57 \simeq 1.8 \times 10^{14}$
- **States:** One queen per column.
 - State space: $8^0 + 8^1 + 8^2 + \dots + 8^8 \simeq 1.9 \times 10^7$
 - Solution space: $8^8 \simeq 1.6 \times 10^7$
- **States:** All possible arrangements of n ($0 \leq n \leq 8$) queens at leftmost n columns with no queen attacked.

Actions: Add a queen to the next column with no queen attacked, or backtrack.

 - State space: 2057.

Tree Search Algorithms



- Offline, simulated exploration of state space by generating successors of already-explored states (a.k.a. **expanding** states)

TREE-SEARCH(*problem*)

```
1  initialize the frontier using the initial state of problem
2  repeat
3      if the frontier is empty
4          return failure
5      choose a leaf node and remove it from the frontier.
6      if the node contains a goal state
7          return the corresponding solution
8      expand the chosen node
9      add the resulting nodes to the frontier
```

Repeated States in Graph Search

- Failure to detect repeated states can turn a linear problem into an exponential one!
- Use a **queue** to record explored states.
- For fast detection of repeated states, **hashing** techniques are usually adopted.



Graph Search Algorithms




GRAPH-SEARCH(*problem*)


```
1  initialize the frontier using the initial state of problem
2  initialize the explored set to be empty
3  repeat
4      if the frontier is empty
5          return failure
6      choose a leaf node and remove it from the frontier.
7      if the node contains a goal state
8          return the corresponding solution
9      add the node to the explored set
10     expand the chosen node
11     if not in the frontier or explored set
12         add the resulting nodes to the frontier
```

Partial Search Tree

(a) The initial state




(b) After expanding Arad




Partial Search Tree

(b) After expanding Arad



(c) After expanding Sibiu




Graph Search, Search Tree, and Frontier Separation

- The frontier **separates** the state space into explored and unexplored regions (**loop invariant proof**).

Tree generated by graph search on Romania map




Separation property of GRAPH-SEARCH



Implementation: States vs. Nodes

- A **state** is a (representation of) a physical configuration
- A **node** is a data structure constituting part of a search tree includes **parent**, **children**, **depth**, **path cost** $g(x)$
- States do not have parents, children, depth, or path cost!



Infrastructure for Search Algorithms

CHILD-NODE(*problem, parent, action*)

return a node *n* with

n.state = *problem.RESULT*(*parent.state, action*)

n.parent = *parent*

n.action = *action*

n.path_cost = *parent.path_cost*

+ *problem.STEP-COST*(*parent.state, action*)

- The appropriate data structure to maintain the frontier is a **queue**.
- Can be
 - FIFO.
 - LIFO (a.k.a. **stack**)
 - **Priority queue**

Tree Search Algorithms


- A **strategy** is defined by picking the **order of node expansion**
- Strategies are evaluated along the following dimensions:
 - **Completeness** - does it always find a solution if one exists?
 - **Optimality** - does it always find a least-cost solution?
 - **Time complexity** - number of nodes generated/expanded
 - **Space complexity** - maximum number of nodes in memory
- Time and space complexity are measured in terms of
 - **b** - **maximum branching factor** of the search tree
 - **d** - **depth** of the least-cost solution
 - **m** - **maximum depth** of the state space (may be ∞)

Uninformed Search Strategies (Blind Search)

- **Uninformed** strategies use only the information available in the problem definition.
- Breadth-first search
- Uniform-cost search
- Depth-first search
- Depth-limited search
- Iterative deepening search

Breadth-First Search (BFS)

- Expand the **shallowest** unexpanded node.
- FIFO queue.



Properties of BFS


- **Completeness:** Yes (if b is finite)
- **Optimality:** No in general; yes when the path cost is a non-decreasing function of the depth of the node.
- **Time complexity:** $1 + b + b^2 + \dots + b^d = O(b^d)$.
Or $O(b^{d+1})$ if goal test is applied after expansion.
- **Space complexity:** $O(b^d)$ (keeps every node in memory)

Space is the big problem; can easily generate nodes at 100MB/sec so
24hrs = 8640GB.

Uniform-Cost Search

- Expand the unexpanded node with the lowest path cost.
- Priority queue ordered by $g(n)$.
- Equivalent to BFS if step costs all equal.

- For TREE-SEARCH, priority queue gives the cheapest path first.
- For GRAPH-SEARCH, if the node is already in the frontier, need to find the minimum cost, and call DECREASEKEY as needed.









Properties of Uniform-Cost Search

- **Completeness:** Yes, if step cost $\geq \epsilon > 0$.
- **Optimality:** Yes - nodes expanded in increasing order of $g(n)$.
- **Time complexity:** # of nodes with $g \leq$ cost of optimal solution.
Maximum depth is given by $1 + \lfloor C^*/\epsilon \rfloor$, where C^* is the cost of the optimal solution.
 $O(b^{1+\lfloor C^*/\epsilon \rfloor})$.
- **Space complexity:** # of nodes with $g \leq$ cost of optimal solution,
 $O(b^{1+\lfloor C^*/\epsilon \rfloor})$.







Depth-First Search (DFS)

- Expand the **deepest** unexpanded node.
- LIFO queue, *i.e.*, put successors at front.



Depth-First Search (DFS)

- Expand the **deepest** unexpanded node.
- LIFO queue, *i.e.*, put successors at front.



Properties of DFS

- **Completeness:** No, fails in infinite-depth spaces, spaces with loops
 - Modify to avoid repeated states along path \rightarrow complete in finite spaces.
- **Optimality:** No.
- **Time complexity:** $O(b^m)$, terrible if m is much greater than d .
 - But if solutions are dense, may be much faster than breadth-first
- **Space complexity:** $O(bm)$, linear space!
 - **Backtracking** technique only generate one successor instead of all successors $\rightarrow O(m)$.



Depth-Limited Search (DLS)



- DFS never terminates if $m \rightarrow \infty$.
- DLS = DFS with depth limit ℓ ,
- Nodes at depth ℓ have no successors
- Recursive implementation:

DEPTH-LIMITED-SEARCH(*problem*, *limit*)

return RECURSIVE-DLS(MAKE-NODE(*problem.initial_state*), *problem*, *limit*)

Depth-Limited Search (DLS)

RECURSIVE-DLS(*node*, *problem*, *limit*)

```
1  if problem.GOAL-TEST(node.state)
2      return SOLUTION(node)
3  elseif limit == 0
4      return cutoff
5  else
6      cutoff_occurred = FALSE
7      for each action in problem.ACTIONS(node.state)
8          child = CHILD-NODE(problem, node, action)
9          result = RECURSIVE-DLS(child, problem, limit - 1)
10         if result == cutoff
11             cutoff_occurred = TRUE
12         elseif result ≠ failure
13             return result
14     if cutoff_occurred
15         return cutoff
16     else
17         return failure
```

Properties of DLS



- **Completeness:** Not complete if $\ell < d$; complete otherwise.
- **Optimality:** Not optimal in general (even if $\ell > d$).
- **Time complexity:** $O(b^\ell)$
- **Space complexity:** $O(b\ell)$, linear space.
- **Two termination conditions:**
 - *failure*: no solution.
 - *cutoff*: no solution within the depth limit.

Iterative-Deepening Search (IDS)



- Call DLS iteratively with **increasing** depth limit.
- Seems to be wasteful, but actually **not**.
- Combine the benefits of BFS and DFS.

ITERATIVE-DEEPENING-SEARCH(*problem*)

```
1  for depth = 0 to  $\infty$ 
2      result = DEPTH-LIMITED-SEARCH(problem, depth)
3      if result  $\neq$  cutoff
4          return result
```


Iterative Deepening Search


Limit = 0



Limit = 1




Limit = 2




Iterative Deepening Search

Limit = 3



Properties of Iterative Deepening Search

- **Completeness:** Yes 
- **Optimality:** No in general; yes when the path cost is a non-decreasing function of the depth of the node.
 - Can be modified to explore uniform-cost trees (optimal).
- **Time complexity:** $db^1 + (d-1)b^2 + \dots + b^d = O(b^d)$
- **Space complexity:** $O(bd)$
- Numerical comparison for $b = 10$, $d = 5$, solution at far right leaf.
 - $N(\text{IDS}) = 50 + 400 + 3,000 + 20,000 + 100,000 = 123,450$
 - $N(\text{BFS}) = 10 + 100 + 1,000 + 10,000 + 100,000 = 111,100$
- Repeated search in IDS is not severe.
- In general, IDS is preferred when **search space is large** and **depth is unknown**.
- We'll talk about more advantages of IDS in **adversarial search**.

Summary of Algorithms

Criterion	BFS	Uniform-Cost	DFS	DLS	IDS
Completeness	Yes ^a	Yes ^b	No	No ^c	Yes ^a
Optimality	No ^d	Yes	No	No	No ^e
Time Complexity	$O(b^d)$	$O(b^{1+\lceil C^*/\epsilon \rceil})$	$O(b^m)$	$O(b^\ell)$	$O(b^d)$
Space Complexity	$O(b^d)$	$O(b^{1+\lceil C^*/\epsilon \rceil})$	$O(bm)$	$O(b\ell)$	$O(bd)$

^aif b is finite

^bif b is finite and step cost $\geq \epsilon$

^cunless $\ell \geq d$

^dunless the path cost is a non-decreasing function of the depth of the node

Summary

- Problem formulation usually requires abstracting away real-world details to define a state space that can feasibly be explored.
 - Initial state.
 - Actions.
 - Transition model.
 - Goal test.
 - Path cost.
- Graph search can be exponentially more efficient than tree search, but usually impractical due to memory requirement.
- Variety of uninformed search strategies judged on the basis of
 - Completeness
 - Optimality
 - Time and space complexity.
- **Iterative deepening** search uses only linear space and not much more time than other uninformed algorithms.